Course Outline

Simon Rayner
simon.rayner@medisin.uio.no

https://github.com/IzmirKCU/IKCU_1

Focus on

- Install Ubuntu for Windows
- Introduction to Python programming
- how to write clean and reusable Python code
- How to debug code and report errors
- How to maintain and collaborate on code
- How to document code
- How to keep an electronic lab book

Course files

You can download from github at

If you are familiar with github, you can pull it from https://github.com/orgs/CBGOUS/repositories

Need java 11. Use SDKMAN to handle java versions (https://sdkman.io/)

```
curl -s "https://get.sdkman.io" | bash
```

Note: you need to restart terminal after installation (or just logout and in again)
or run

```
source "$HOME/.sdkman/bin/sdkman-init.sh”
```

To install Corretto 17:

```
sdk install java 17.0.6-amzn
```

To install Corretto 8:

```
sdk install java 8.0.442-amzn
```

hairpin.fa

From mirbase.org, a list of all hairpin sequences, for all species

Question 1

How many sequences in the file?

Question 2

How many species in the file?

Question 3

How many human sequences in the file?

Install Ubuntu on Windows

Install Ubuntu on Windows

Install Python

Install Ubuntu on Windows

Install Java (sdkman)

Install Java (sdkman)

Need java 11. Use SDKMAN to handle java versions (https://sdkman.io/)

```
curl -s "https://get.sdkman.io" | bash
```

Note: you need to restart terminal after installation (or just logout and in again)
or run

```
source "$HOME/.sdkman/bin/sdkman-init.sh”
```

To install Corretto 17:

```
sdk install java 17.0.6-amzn
```

To install Corretto 8:

```
sdk install java 8.0.442-amzn
```

Some basic Linux commands

pwd
ls
cd
wc –l
grep
sed
awk
find
xargs

piping commands using the ‘|’ character

Some basic Linux commands

Install Java (sdkman)

GC Calculation

Find the GC content of the fasta file `hairpin.fa`

```
cat hairpin.fa|more
```

```
>cel-let-7 MI0000001 Caenorhabditis elegans let-7 stem-loop
UACACUGUGGAUCCGGUGAGGUAGUAGGUUGUAUAGUUUGGAAUAUUACCACCGGUGAAC
UAUGCAAUUUUCUACCUUACCGGAGACAGAACUCUUCGA
>cel-lin-4 MI0000002 Caenorhabditis elegans lin-4 stem-loop
AUGCUUCCGGCCUGUUCCCUGAGACCUCAAGUGUGAGUGUACUAUUGAUGCUUCACACCU
GGGCUCUCCGGGUACCAGGACGGUUUGAGCAGAU
>cel-mir-1 MI0000003 Caenorhabditis elegans miR-1 stem-loop
AAAGUGACCGUACCGAGCUGCAUACUUCCUUACAUGCCCAUACUAUAUCAUAAAUGGAUA
UGGAAUGUAAAGAAGUAUGUAGAACGGGGUGGUAGU
>cel-mir-2 MI0000004 Caenorhabditis elegans miR-2 stem-loop
UAAACAGUAUACAGAAAGCCAUCAAAGCGGUGGUUGAUGUGUUGCAAAUUAUGACUUUCA
UAUCACAGCCAGCUUUGAUGUGCUGCCUGUUGCACUGU
```

Let’s modify the file so that we have one sequence/line

```
awk '/^>/ {printf("\n%s\n",$0);next; } { printf("%s",$0);} END {printf("\n");}’ \
    < hairpin.fa > hairpin_flat.fa
```

```
cat hairpin_flat.fa|more
```

```
>cel-let-7 MI0000001 Caenorhabditis elegans let-7 stem-loop
UACACUGUGGAUCCGGUGAGGUAGUAGGUUGUAUAGUUUGGAAUAUUACCACCGGUGAACUAUGCAAUUUUCUACCUUACCGGAGACAGAACUCUUCGA
>cel-lin-4 MI0000002 Caenorhabditis elegans lin-4 stem-loop
AUGCUUCCGGCCUGUUCCCUGAGACCUCAAGUGUGAGUGUACUAUUGAUGCUUCACACCUGGGCUCUCCGGGUACCAGGACGGUUUGAGCAGAU
>cel-mir-1 MI0000003 Caenorhabditis elegans miR-1 stem-loop
AAAGUGACCGUACCGAGCUGCAUACUUCCUUACAUGCCCAUACUAUAUCAUAAAUGGAUAUGGAAUGUAAAGAAGUAUGUAGAACGGGGUGGUAGU
>cel-mir-2 MI0000004 Caenorhabditis elegans miR-2 stem-loop
UAAACAGUAUACAGAAAGCCAUCAAAGCGGUGGUUGAUGUGUUGCAAAUUAUGACUUUCAUAUCACAGCCAGCUUUGAUGUGCUGCCUGUUGCACUGU
```

Let’s modify the file so that we have one sequence/line

```
awk '/^>/ {printf("\n%s\n",$0);next; } { printf("%s",$0);} END {printf("\n");}’ < hairpin.fa
```

We are going to work with the code and data in the folder
../BINF_M612/day1/GCCalculation/software

This folder contains two java files CalcGC.jar and GCCalc.jar

They both calculate average GC content for an input file of fasta sequences

We can find out how to run the program by typing

```
$ java -jar day1/GCCalculation/software/GCcalc.jar -h

GCCalc
initializing
parse arguments
  ==================================================================\
  | GCCalc  :                                                       \
  |   Java code to calculate GC percentage in a FastA file          \
  ==================================================================\
usage: command line options
 -f,--sequence file <arg>   sequence file in FASTA format
 -h,--help                  view help
```

So, we just need to specify a fasta file. Let's start with the file data/GCtest.fa

(This corresponds to an alignment of sequences for the 3'UTR of the **Lysine Methyltransferase 5B** gene)

We've tried with a small test dataset. Let's repeat with a real one.

```
$ java -jar day1/GCCalculation/software/GCCalc.jar
-f day1/GCCalculation/data/GCtest.fa
```

```
GCCalc
initializing
parse arguments
fasta input file is <day1/GCCalculation/data/ENSG00000110066___ENST00000441488___2___KMT5B__uniq_aln.fa>
read <16> sequences from file
average GC value of all sequences is <43.35%>
```

And repeat using CalcGC.jar

```
$ java -jar day1/GCCalculation/software/CalcGC.jar
-f day1/GCCalculation/data/GCtest.fa
```

```
CalcGC
initializing
parse arguments
fasta input file is <day1/GCCalculation/data/ENSG00000110066___ENST00000441488___2___KMT5B__uniq_aln.fa>
read <16> sequences from file
average GC value of all sequences is <43.35%>
```

So, two programs give the same results, which is encouraging

We've checked the programs using a simple test dataset.
Now let's try running the programs against the hairpin.fa file

What do you find?

How can you figure out what is going on?

Data set is very large – so let’s use a simpler test set

Programming in Python

We started with calculating the average GC content of all the sequences in the hairpin.fa file using two different Java programs

We installed Java using SDKMAN
This allows us to run different versions of Java, which is quite handy if you are running programs you have downloaded from other sites

nextflow

Documentation Training Forums Examples Tools Resources

Reproducible Scientific Workflows at Scale

Nextflow enables scalable, reproducible, and portable scientific workflows for research and production use cases.

Documentation | Community forum

Check prerequisites
Java 17 or later is required

Make sure Java 17 or later is installed on your computer by using the command:

```
java -version
```

Note: If you are having trouble installing or upgrading Java check out our documentation here.

Step 1 | Step 2 | Step 3

For example, Nextflow requires Java 17

But Picard, another popular tool only requires Java 8

Picard

build passing

A set of command line tools (in Java) for manipulating high-throughput sequencing (HTS) data and formats such as SAM/BAM/CRAM and VCF.

Latest Jar Release | Source Code ZIP File | Source Code TAR Ball | View On GitHub

Picard is a set of command line tools for manipulating high-throughput sequencing (HTS) data and formats such as SAM/BAM/CRAM and VCF. These file formats are defined in the **Hts-specs** repository. See especially the **SAM specification** and the **VCF specification**.

For the tools to run properly, you must have Java 1.8 installed. To check your java version by open your terminal application and run the following command:

```
java -version
```

If the output looks something like `java version "1.8.x"`, you are good to go. If not, you may need to update your version; see the **Oracle Java website** to download the latest JDK.

We found that running the two different Java programs sometimes returned different values for GC %

We didn’t have the source code, but we found the error by creating a simple test dataset to put into the programs
i.e., rather than calculating the GC% for 38000 sequences, we created a test file containing the sequence

```
>test1
AACCGGTT
```

Now we are going to do the same thing by writing some Python code

We will do this by starting with a simple Python program and run in two different ways

First of all, we will run it from inside a terminal window
Then we will run it inside Jupyter

Programming in Python

Writing code inside Jupyter Notebook

Open a command window
(in Windows you can do this by typing
`<CTRL>+<SHIFT>+P`)

When a window opens, type

```
jupyter notebook
```

After a while (depending on your computer's speed), a webpage will open that looks something like this

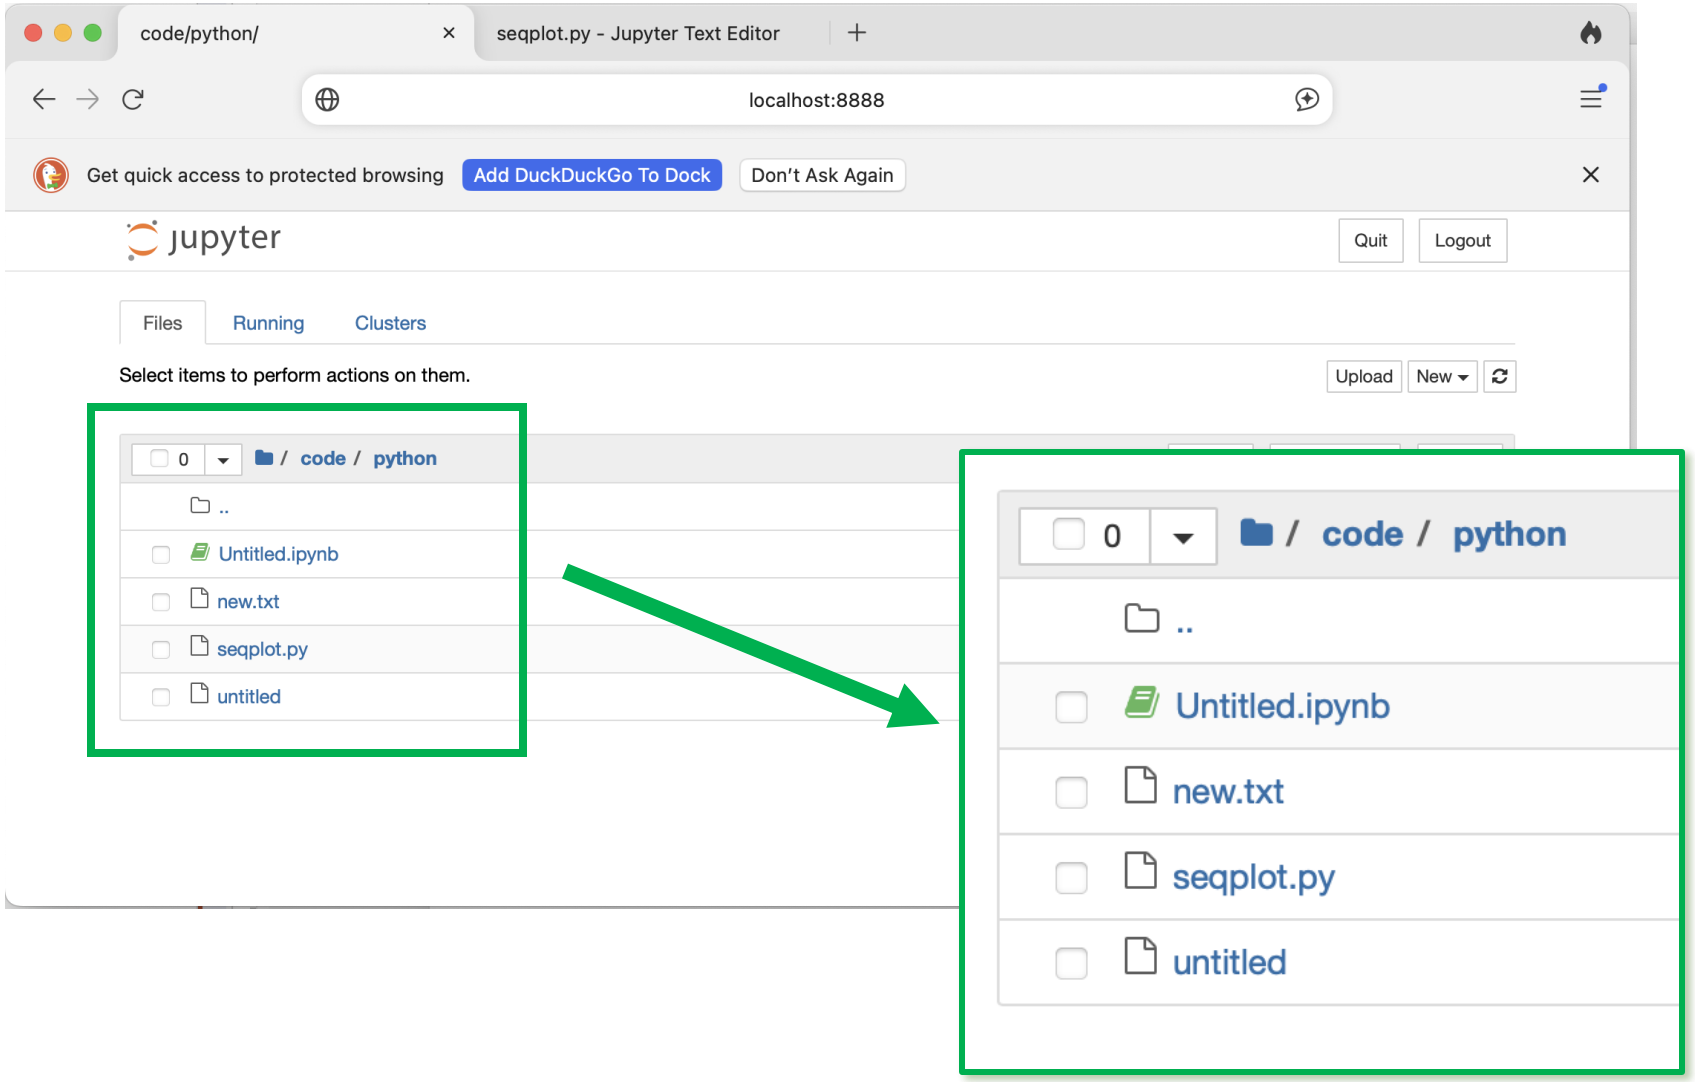

It won't look exactly the same, because you need to move to the folder where you downloaded the code

For example, before i started **jupyter**, i changed the directory to the folder where i downloaded the code

```
(base) MacBook-Pro-2:IKCU2 simonray$ pwd
/Users/simonray/uio_dropbox_sr/myTeaching/IKCU2
(base) MacBook-Pro-2:IKCU2 simonray$ tree
.
├── ~$ikcu_2025_outline.pptx
├── code
│   ├── java
│   └── python
│       ├── new.txt
│       ├── seqplot.py
│       ├── untitled
│       └── Untitled.ipynb
├── data
│   ├── GCtest.fa
│   ├── hairpin_flat.fa
│   └── hairpin.fa
└── ikcu_2025_outline.pptx

4 directories, 9 files
(base) MacBook-Pro-2:IKCU2 simonray$
```

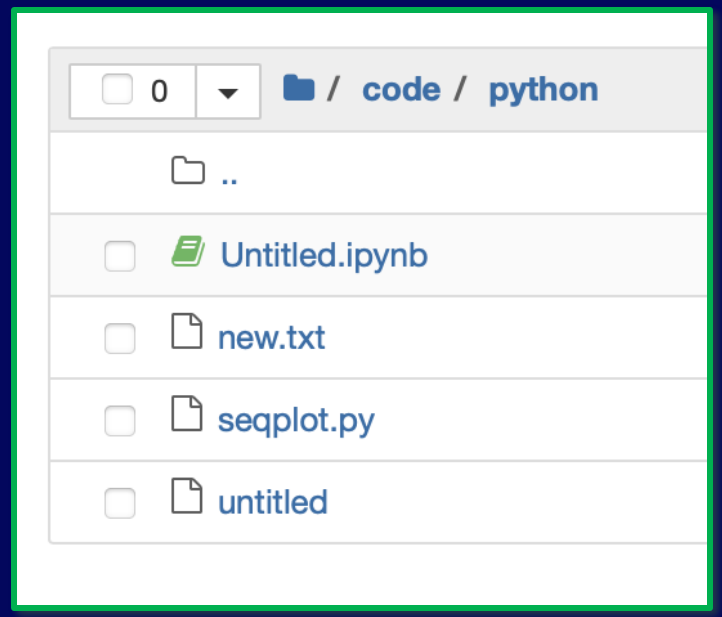

Programming in Python

Running code inside a virtual environment

How would you describe a bicycle?

- Two wheels
- Handlebar
- Saddle
- frame

Here is a bike

Here is another bike

The parts are incompatible (e.g., you can’t change the wheels

This is a bit like the problem you face with Python

```
import sklearn
import tensorflow as tf
```

They both use Pandas, but may require different versions of the Pandas package

What about Anaconda or MiniConda?

Anaconda gives you whatever version happens to be in the package. So it tries to make sure the are no dependency issues, but it may break other programs you already have installed

Also, Anaconda gives you many other packages you may never use







Virtual Environments

Virtual Enviroments give you a way to create a custom environment to run your code

VE1

VE 2

Reproducible Research

miRBASE ANNOTATION : HIGH CONFIDENCE SETS

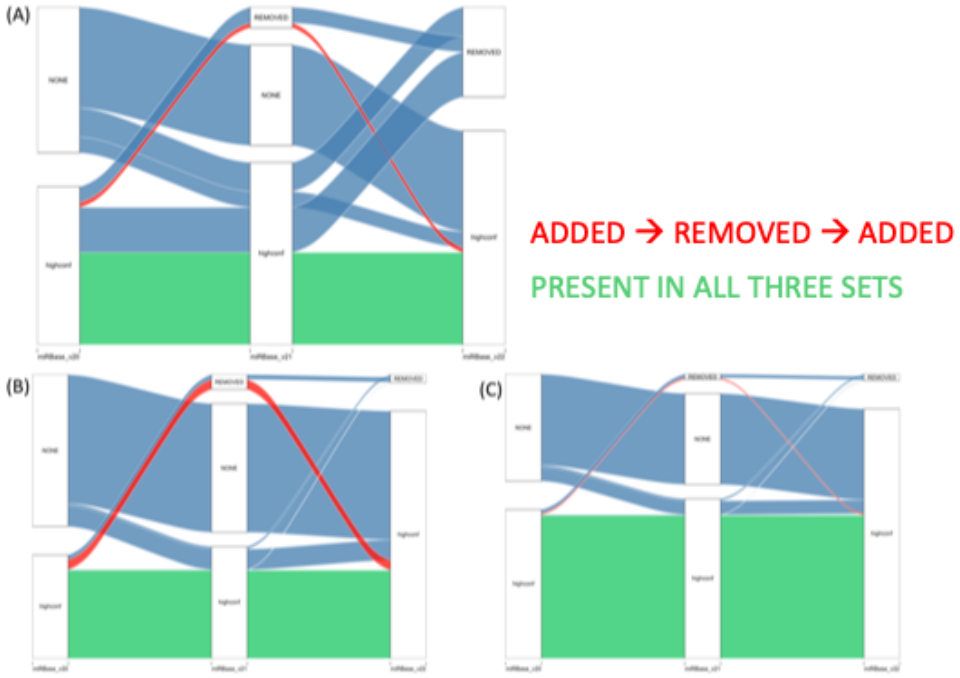

TO ADDRESS SOME OF THESE ANNOTATION PROBLEMS miRBASE RELEASED A HIGH CONFIDENCE ANNOTATION SET

3298 DISTINCT HAIRPIN PRECURSORS ACROSS THE THREE RELEASES,
ONLY 925 (231 HUMAN HAIRPINS) ARE PRESENT ACROSS ALL THREE RELEASES

We saw that the high confidence miRBase set changed between each version

How does the variation among different miRBase releases affect an analysis?

Variation in miRBase Annotation

miRBase: Growth

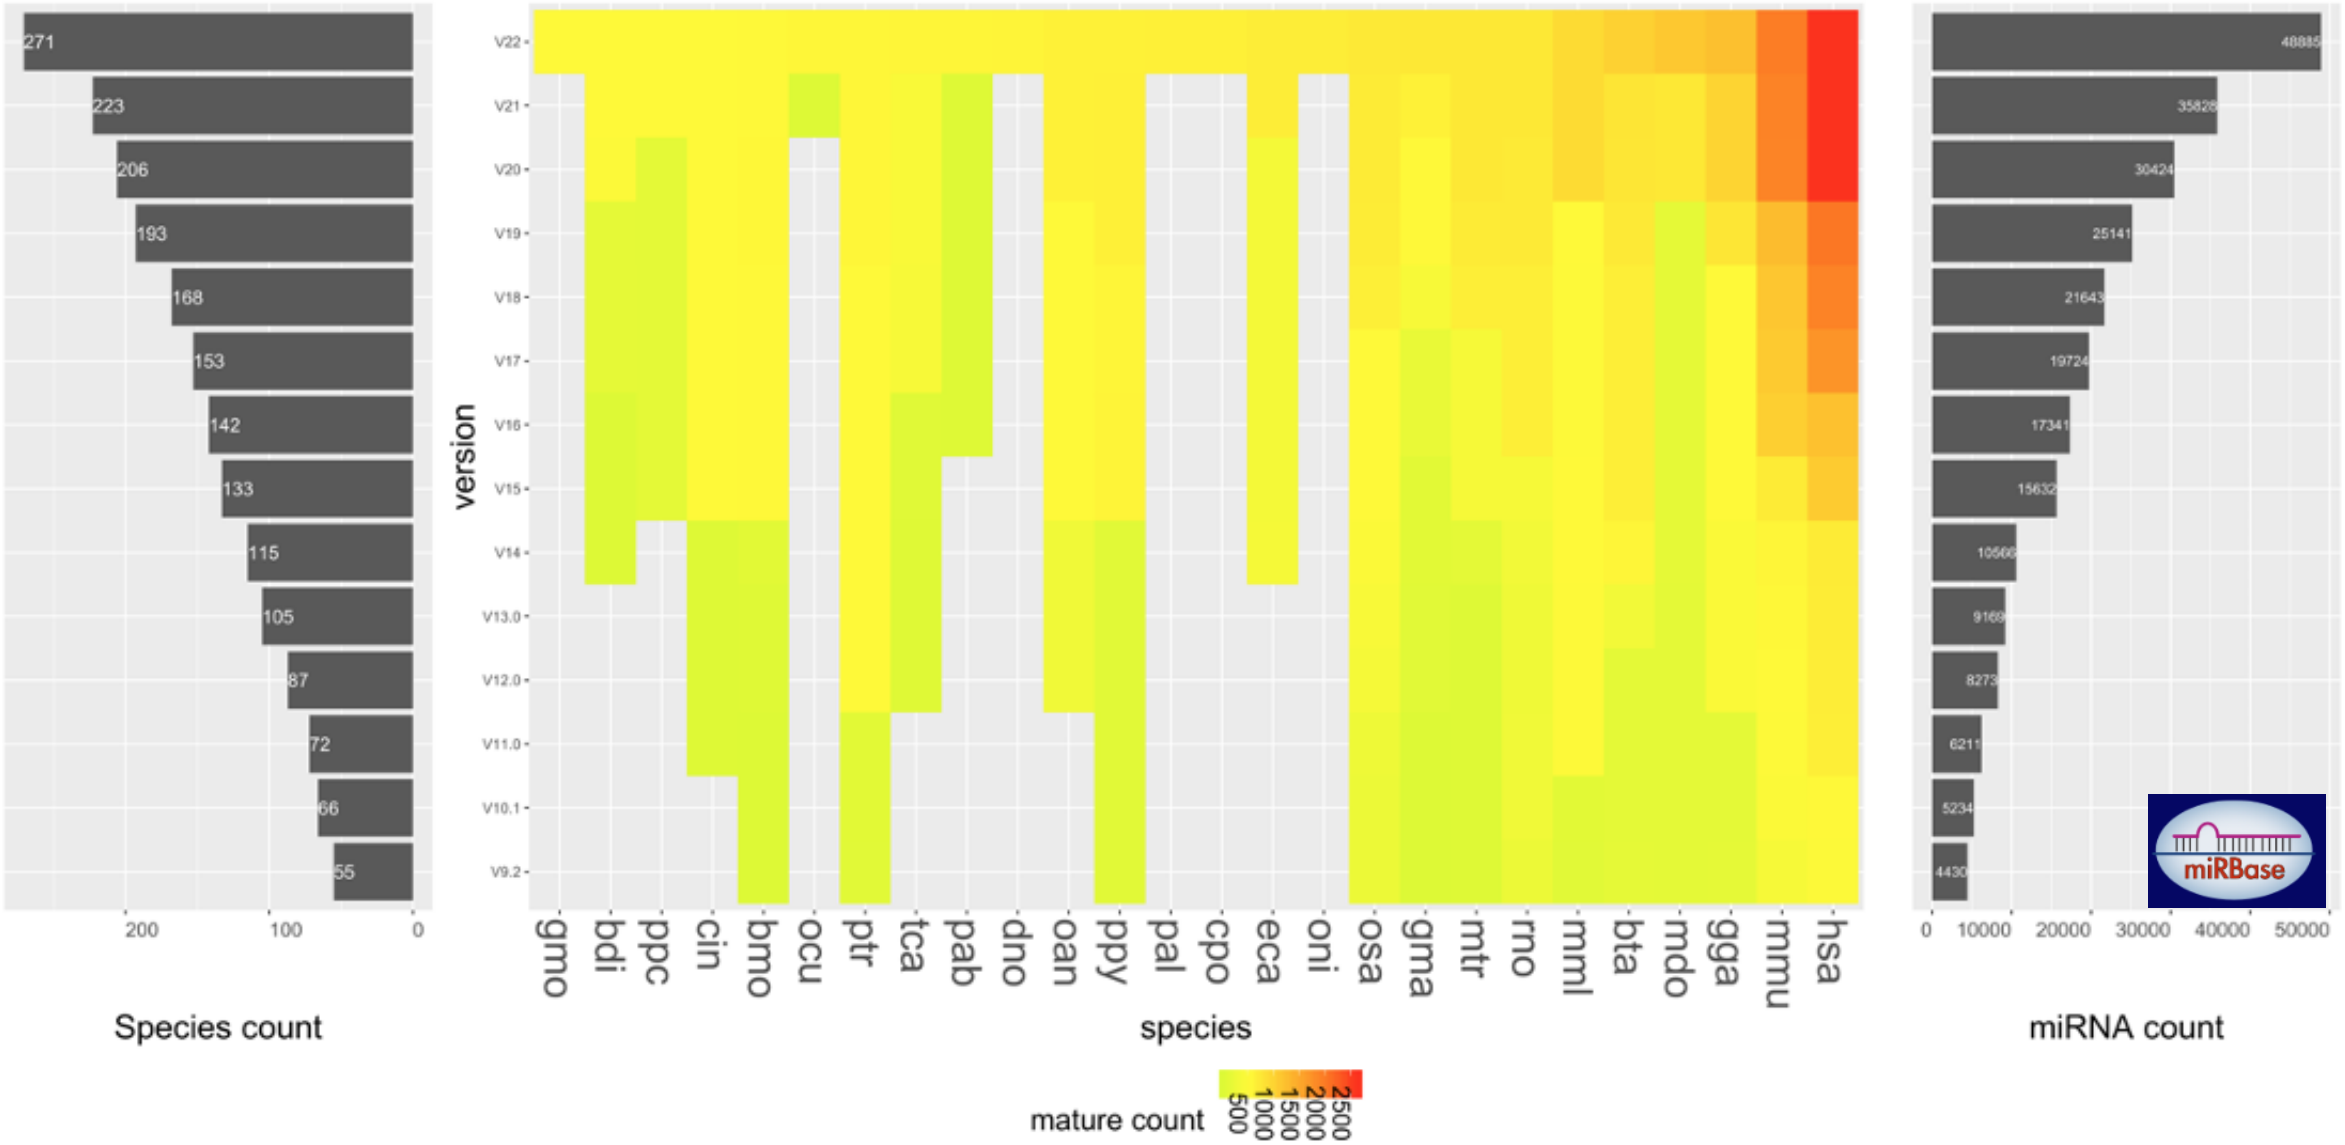

We are going to look at the impact of using different annotation and parameter values on the analysis results

To do this, we are going to run the mapping part of a NGS analysis using smallRNA seq data using the `bowtie` mapping tool

To do the analysis, we need some reads, and a reference genome
It will take too long to work with a whole NGS dataset + Reference Genome, so we are going to use a test dataset, `smallRNA_reads.fa`
And a shorter reference genome containing only `chr 10` and `chr X`.

Osteoarthritis and Cartilage Open 3 (2021) 100131

Contents lists available at ScienceDirect

Osteoarthritis and Cartilage Open

journal homepage: www.elsevier.com/journals/osteoarthritis-and-cartilage-open/2665-9131

Experimental Protocol

A bioinformatics approach to microRNA-sequencing analysis

Pratibha Potla[a,b,1], Shabana Amanda Ali[c,**,1], Mohit Kapoor[a,b,d,*]

also the most correct reads based on their UMI tag. Since it is imperative to retain and trim only those reads having the 3' adapter the sequencer will read into the adapter in order to capture miRNAs. The utility of UMIs is only seen post-alignment as there is greater confidence in the genomic read location. Since the length of mature miRNAs is known to be around 22–25 bp, the final raw read filtering step is to trim the reads to retain only the expected miRNA read lengths with some leniency, to remove reads that are either too short (<18 bp) and too long (>30 bp). The result of UMI analysis and read filtering is a set of good quality raw sequences, ready to be processed for any analysis, such as alignment.

RNA BIOLOGY
2019, VOL. 16, NO. 11, 1534–1546
https://doi.org/10.1080/15476286.2019.1637680

RESEARCH PAPER

miRBaseMiner, a tool for investigating miRBase content

Xiangfu Zhong, Fatima Heinicke, and Simon Rayner*

Department of Medical Genetics, Oslo University Hospital and University of Oslo, Oslo, Norway

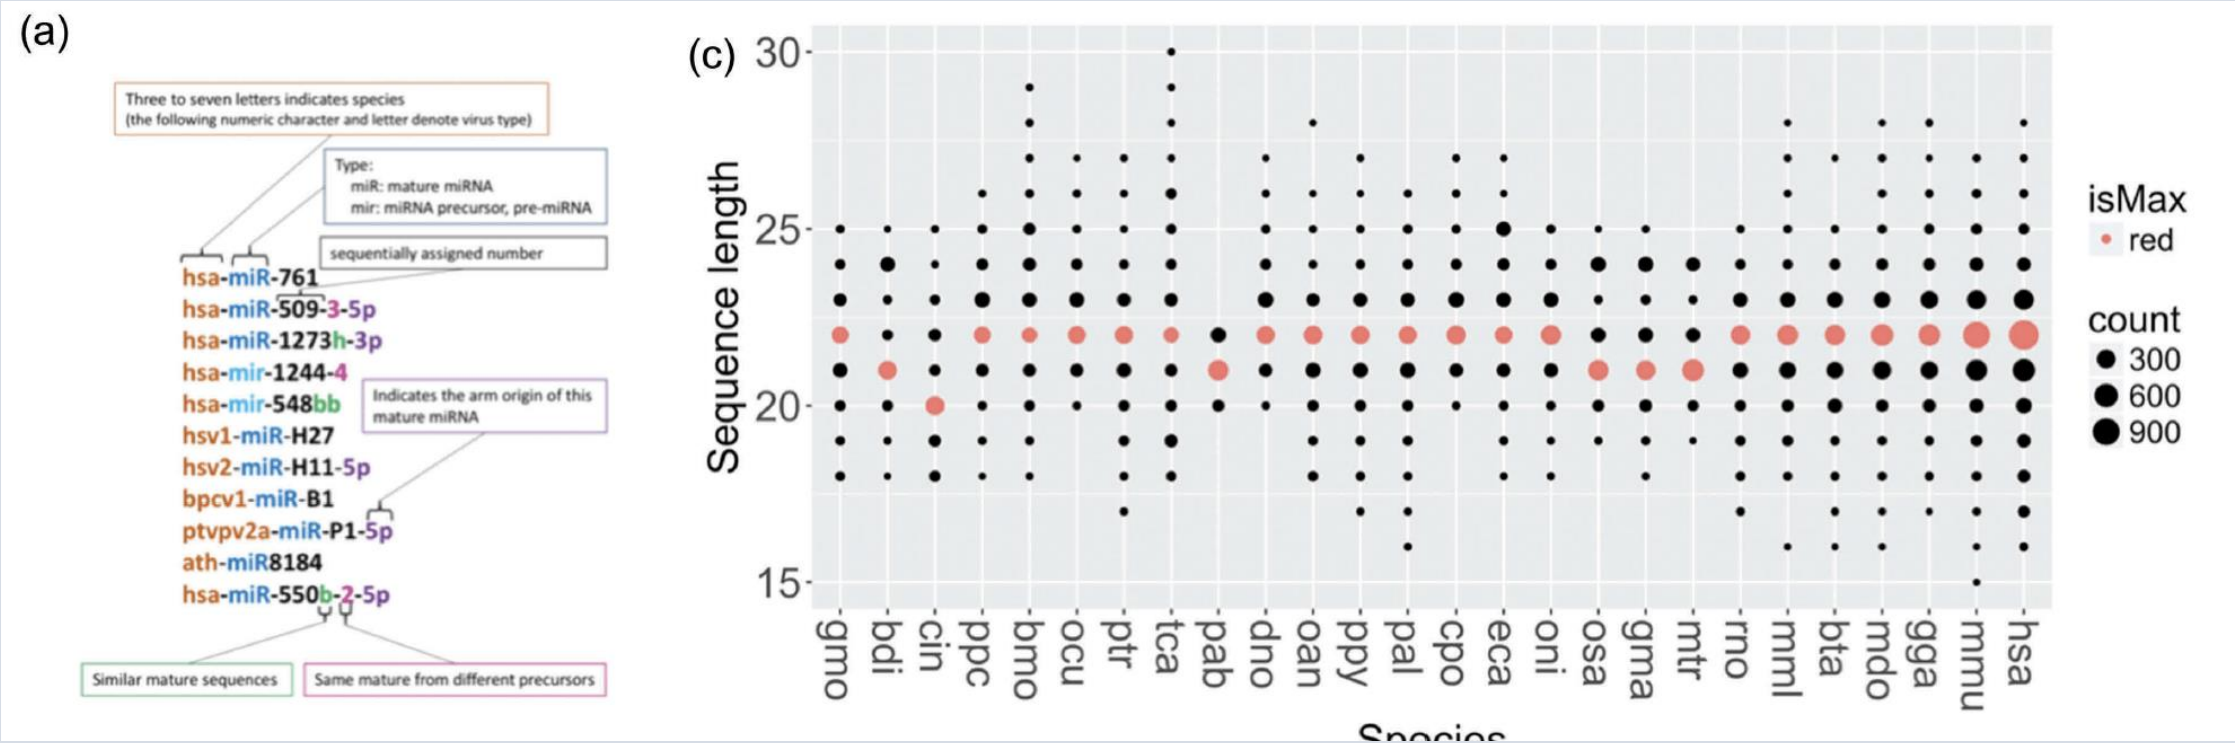

What is the length distribution of human miRNAs?

RNA BIOLOGY
2020, VOL. 17, NO. 1, 75–86
https://doi.org/10.1080/15476286.2019.1667741

Taylor & Francis
Taylor & Francis Group

RESEARCH PAPER

Systematic assessment of commercially available low-input miRNA library preparation kits

Fatima Heinicke[a], Xiangfu Zhong[a], Manuela Zucknick[b], Johannes Breidenbach[c], Arvind Y. M. Sundaram[a], Siri T. Flåm[a], Magnus Leithaug[a], Marianne Dalland[a], Andrew Farmer[d], Jordana M. Henderson[e], Melanie A. Hussong[f], Pamela Moll[g], Loan Nguyen[h], Amanda McNulty[d], Jonathan M. Shaffer[f], Sabrina Shore[e**], Hoichong Karen Yip[h], Jana Vitkovska[g], Simon Rayner[a], Benedicte A Lie[a*], and Gregor D. Gilfillan[a*]

Bioinformatic analysis

Read mapping and reference sequences

Primary base calling and quality scoring was performed using RTA v1.18.66.4 (Illumina), followed by demultiplexing and processing with Bcl2fastq v2.18.0.12 (Illumina).

For trimming of the 3′ adapter, we followed adapter trimming instructions according to each manufacturer (cutadapt v1.15[41] with parameter –m 10 was used in all cases). Detailed information about adapter sequences is provided in the Supplementary Material and Methods.

Read mapping was performed using bowtie v1.1.2[42] with parameters –a and –norc. No mismatch was allowed. As

```
bowtie -x bowtie/hsa_chr10 -f ngsdata/smallRNA_reads.fa -n 0
```

We are going to look at the impact of using different annotation and parameter values on the analysis results

To do this, we are going to run the mapping part of a NGS analysis using smallRNA seq data using the `bowtie` mapping tool

To do the analysis, we need some reads, and a reference genome
It will take too long to work with a whole NGS dataset + Reference Genome, so we are going to use a test dataset, `smallRNA_reads.fa`
And a shorter reference genome containing only `chr 10` and `chr X`.

Let's see how this affects read mapping ....

We need the following software

bowtie https://bowtie-bio.sourceforge.net/index.shtml

samtools
bedtools
bgzip

Installed using apt-get install

1. Map the reads to the reference genome

```
bowtie -x bowtie/hsa_chr10 -f ngsdata/smallRNA_reads.fa
bowtie -x bowtie/hsa_chr10 -f ngsdata/smallRNA_reads.fa -S 10.sam
```

2. Map the reads to the reference genome and write to 10.sam

```
samtools view -bo 10.bam 10.sam
```

3. Convert the alignment results to binary format (less space and faster to process)

```
bedtools intersect -a mirbase/21/hsa_s.gff3 -b 10.bam
```

4. Find which reads overlap the features in the gff file

```
chr11	.	miRNA_primary_transcript	2134134	2134209	.	-	.	ID=MI0002467;Alias=MI0002467;Name=hsa-mir-483
chr11	.	miRNA	2134181	2134202	.	-	.	ID=MIMAT0004761;Alias=MIMAT0004761;Name=hsa-miR-483-5p;Derives_from=MI0002467
chr11	.	miRNA	2134142	2134162	.	-	.	ID=MIMAT0002173;Alias=MIMAT0002173;Name=hsa-miR-483-3p;Derives_from=MI0002467
```

```
chr11	.	miRNA_primary_transcript	2134134	2134209	.
-
chr11	.	miRNA	2134181	2134202	.	-	.

chr11	.	miRNA	2134142	2134162	.	-	.
```

IN REALITY, THERE CAN BE MULTIPLE ISOFORMS (isomiRs) OF AN miRNA

2134142 2134162

```
+-----------------------------------------------------------------------------+
+GAGGGGGAAGACGGGAGGAAAGAAGGGAGUGGUUCCAUCACGCCUCCUCACUCCUCUCCUCCCGUCUUCUCCUCUC+
       GAAGACGGGAGGAAATAAGGGAG ✗
     GGAAGACGGGAGGAAAGAAGGGAG ✗
       GAAAGACGGGAGGAAAGAAGGG ✗
       GAAAGACGGGAGGAAAGAAGGGG ✓
```

COMMONLY, BUT NOT ALWAYS, ONLY THE MATURE miRNA IS COUNTED

THE PRESENCE OF isomiRs COMPLICATES THE TARGETING PROCESS (SHIFTED SEED REGION)